

Logistic Regression

October 15, 2024

1 Logistic Regression

```
[13]: # Importing the libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import os
```

```
[14]: # Importing the dataset
os.chdir("C:\\Users\\ddaya\\OneDrive\\Documents\\Python_programming")
dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
```

```
[15]: # Splitting the dataset into the Training set and Test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25,
    ↪random_state = 0)
print(X_train)
print(y_train)
print(X_test)
print(y_test)
```

```
[[ 37 71000]
 [ 36 50000]
 [ 48 29000]
 [ 30 87000]
 [ 32 18000]
 [ 32 100000]
 [ 47 25000]
 [ 40 75000]
 [ 29 47000]
 [ 27 17000]
 [ 31 76000]
 [ 48 41000]
 [ 26 35000]
 [ 31 15000]
 [ 41 51000]
```

[24 27000]
[29 75000]
[38 61000]
[47 30000]
[20 74000]
[33 31000]
[35 25000]
[35 75000]
[24 55000]
[23 66000]
[26 43000]
[32 117000]
[29 83000]
[35 44000]
[30 15000]
[26 80000]
[35 88000]
[33 113000]
[18 86000]
[26 86000]
[28 59000]
[22 81000]
[20 82000]
[33 69000]
[24 32000]
[37 55000]
[46 59000]
[59 83000]
[41 59000]
[37 72000]
[31 68000]
[31 89000]
[30 135000]
[31 118000]
[32 18000]
[31 58000]
[45 22000]
[18 68000]
[45 22000]
[29 80000]
[28 87000]
[29 83000]
[18 52000]
[27 57000]
[21 88000]
[26 84000]
[27 84000]
[22 63000]

[34 112000]
[30 62000]
[28 55000]
[20 86000]
[21 68000]
[27 96000]
[18 82000]
[19 19000]
[30 116000]
[41 30000]
[25 79000]
[26 52000]
[32 86000]
[35 27000]
[28 79000]
[33 51000]
[35 71000]
[24 58000]
[35 20000]
[36 75000]
[19 25000]
[32 135000]
[35 108000]
[21 72000]
[26 15000]
[45 26000]
[26 72000]
[35 23000]
[23 63000]
[30 17000]
[27 90000]
[29 43000]
[42 80000]
[27 137000]
[30 89000]
[26 32000]
[41 45000]
[21 16000]
[35 59000]
[31 18000]
[28 37000]
[20 49000]
[26 17000]
[40 57000]
[29 61000]
[22 55000]
[20 23000]
[22 27000]

4

```

[ 28 59000]
[ 39 71000]
[ 26 30000]
[ 49 28000]
[ 23 48000]
[ 42 65000]
[ 23 20000]
[ 25 33000]
[ 28 84000]
[ 30 80000]
[ 27 89000]
[ 35 73000]
[ 29 148000]
[ 46 23000]
[ 31 74000]
[ 26 15000]]
[0 0 1 0 0 0 1 0 0 0 1 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0
 0 0 1 1 0 0]

```

```

[16]: # Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
print(X_train)
print(X_test)

```

```

[[ 8.18877097e-01  3.33231348e-01]
 [ 6.91518326e-01 -3.63289541e-01]
 [ 2.21982358e+00 -1.05981043e+00]
 [-7.26342990e-02  8.63913930e-01]
 [ 1.82083243e-01 -1.42465470e+00]
 [ 1.82083243e-01  1.29509353e+00]
 [ 2.09246481e+00 -1.19248108e+00]
 [ 1.20095341e+00  4.65901993e-01]
 [-1.99993070e-01 -4.62792525e-01]
 [-4.54710612e-01 -1.45782237e+00]
 [ 5.47244718e-02  4.99069655e-01]
 [ 2.21982358e+00 -6.61798493e-01]
 [-5.82069382e-01 -8.60804461e-01]
 [ 5.47244718e-02 -1.52415769e+00]
 [ 1.32831218e+00 -3.30121880e-01]
 [-8.36786924e-01 -1.12614575e+00]
 [-1.99993070e-01  4.65901993e-01]
 [ 9.46235868e-01  1.55473413e-03]
 [ 2.09246481e+00 -1.02664277e+00]
 [-1.34622201e+00  4.32734332e-01]
 [ 3.09442014e-01 -9.93475107e-01]

```

[5.64159555e-01 -1.19248108e+00]
 [5.64159555e-01 4.65901993e-01]
 [-8.36786924e-01 -1.97451234e-01]
 [-9.64145695e-01 1.67393041e-01]
 [-5.82069382e-01 -5.95463170e-01]
 [1.82083243e-01 1.85894377e+00]
 [-1.99993070e-01 7.31243284e-01]
 [5.64159555e-01 -5.62295509e-01]
 [-7.26342990e-02 -1.52415769e+00]
 [-5.82069382e-01 6.31740300e-01]
 [5.64159555e-01 8.97081591e-01]
 [3.09442014e-01 1.72627313e+00]
 [-1.60093955e+00 8.30746268e-01]
 [-5.82069382e-01 8.30746268e-01]
 [-3.27351841e-01 -6.47805886e-02]
 [-1.09150447e+00 6.64907961e-01]
 [-1.34622201e+00 6.98075623e-01]
 [3.09442014e-01 2.66896025e-01]
 [-8.36786924e-01 -9.60307445e-01]
 [8.18877097e-01 -1.97451234e-01]
 [1.96510603e+00 -6.47805886e-02]
 [3.62077006e+00 7.31243284e-01]
 [1.32831218e+00 -6.47805886e-02]
 [8.18877097e-01 3.66399009e-01]
 [5.47244718e-02 2.33728364e-01]
 [5.47244718e-02 9.30249252e-01]
 [-7.26342990e-02 2.45596168e+00]
 [5.47244718e-02 1.89211143e+00]
 [1.82083243e-01 -1.42465470e+00]
 [5.47244718e-02 -9.79482500e-02]
 [1.83774726e+00 -1.29198406e+00]
 [-1.60093955e+00 2.33728364e-01]
 [1.83774726e+00 -1.29198406e+00]
 [-1.99993070e-01 6.31740300e-01]
 [-3.27351841e-01 8.63913930e-01]
 [-1.99993070e-01 7.31243284e-01]
 [-1.60093955e+00 -2.96954218e-01]
 [-4.54710612e-01 -1.31115911e-01]
 [-1.21886324e+00 8.97081591e-01]
 [-5.82069382e-01 7.64410946e-01]
 [-4.54710612e-01 7.64410946e-01]
 [-1.09150447e+00 6.78900569e-02]
 [4.36800784e-01 1.69310546e+00]
 [-7.26342990e-02 3.47223955e-02]
 [-3.27351841e-01 -1.97451234e-01]
 [-1.34622201e+00 8.30746268e-01]
 [-1.21886324e+00 2.33728364e-01]
 [-4.54710612e-01 1.16242288e+00]

[-1.60093955e+00 6.98075623e-01]
 [-1.47358078e+00 -1.39148704e+00]
 [-7.26342990e-02 1.82577611e+00]
 [1.32831218e+00 -1.02664277e+00]
 [-7.09428153e-01 5.98572639e-01]
 [-5.82069382e-01 -2.96954218e-01]
 [1.82083243e-01 8.30746268e-01]
 [5.64159555e-01 -1.12614575e+00]
 [-3.27351841e-01 5.98572639e-01]
 [3.09442014e-01 -3.30121880e-01]
 [5.64159555e-01 3.33231348e-01]
 [-8.36786924e-01 -9.79482500e-02]
 [5.64159555e-01 -1.35831938e+00]
 [6.91518326e-01 4.65901993e-01]
 [-1.47358078e+00 -1.19248108e+00]
 [1.82083243e-01 2.45596168e+00]
 [5.64159555e-01 1.56043482e+00]
 [-1.21886324e+00 3.66399009e-01]
 [-5.82069382e-01 -1.52415769e+00]
 [1.83774726e+00 -1.15931341e+00]
 [-5.82069382e-01 3.66399009e-01]
 [5.64159555e-01 -1.25881640e+00]
 [-9.64145695e-01 6.78900569e-02]
 [-7.26342990e-02 -1.45782237e+00]
 [-4.54710612e-01 9.63416914e-01]
 [-1.99993070e-01 -5.95463170e-01]
 [1.45567095e+00 6.31740300e-01]
 [-4.54710612e-01 2.52229700e+00]
 [-7.26342990e-02 9.30249252e-01]
 [-5.82069382e-01 -9.60307445e-01]
 [1.32831218e+00 -5.29127848e-01]
 [-1.21886324e+00 -1.49099003e+00]
 [5.64159555e-01 -6.47805886e-02]
 [5.47244718e-02 -1.42465470e+00]
 [-3.27351841e-01 -7.94469139e-01]
 [-1.34622201e+00 -3.96457202e-01]
 [-5.82069382e-01 -1.45782237e+00]
 [1.20095341e+00 -1.31115911e-01]
 [-1.99993070e-01 1.55473413e-03]
 [-1.09150447e+00 -1.97451234e-01]
 [-1.34622201e+00 -1.25881640e+00]
 [-1.09150447e+00 -1.12614575e+00]
 [2.09246481e+00 -1.35831938e+00]
 [1.07359464e+00 -6.28630832e-01]
 [-7.09428153e-01 9.63416914e-01]
 [5.64159555e-01 -7.61301477e-01]
 [-4.54710612e-01 -9.93475107e-01]
 [-1.09150447e+00 -1.42465470e+00]

```

[-1.47358078e+00  7.97578607e-01]
[-5.82069382e-01  6.64907961e-01]
[-7.09428153e-01  6.31740300e-01]
[-3.27351841e-01  7.97578607e-01]
[ 3.09442014e-01 -1.09297809e+00]
[ 2.09246481e+00 -3.96457202e-01]
[ 5.64159555e-01  1.34225380e-01]
[ 3.09442014e-01  2.92030893e+00]
[-9.64145695e-01  6.98075623e-01]
[ 6.91518326e-01 -2.96954218e-01]
[-4.54710612e-01 -2.30618895e-01]]
[[-4.54710612e-01 -9.79482500e-02]
[-8.36786924e-01  9.30249252e-01]
[ 1.82083243e-01  1.95844675e+00]
[-8.36786924e-01 -1.97451234e-01]
[ 1.07359464e+00  1.55473413e-03]
[-4.54710612e-01  8.97081591e-01]
[ 1.82083243e-01  2.95347660e+00]
[ 1.20095341e+00 -4.62792525e-01]
[-1.47358078e+00 -1.32515172e+00]
[ 5.64159555e-01 -3.63289541e-01]
[-7.26342990e-02  1.52726716e+00]
[ 4.36800784e-01 -1.19248108e+00]
[-4.54710612e-01 -9.79482500e-02]
[-3.27351841e-01  2.05794974e+00]
[ 8.18877097e-01 -9.27139784e-01]
[-1.99993070e-01 -1.09297809e+00]
[-3.27351841e-01 -5.62295509e-01]
[ 1.96510603e+00 -1.09297809e+00]
[-7.09428153e-01  8.63913930e-01]
[-1.60093955e+00 -5.62295509e-01]
[-8.36786924e-01 -1.39148704e+00]
[-7.26342990e-02 -3.96457202e-01]
[-1.47358078e+00  4.99069655e-01]
[ 1.20095341e+00 -6.47805886e-02]
[-4.54710612e-01 -1.35831938e+00]
[ 9.46235868e-01  6.31740300e-01]
[ 5.64159555e-01 -2.63786557e-01]
[-3.27351841e-01 -6.47805886e-02]
[ 1.07359464e+00  3.33231348e-01]
[-5.82069382e-01 -1.02664277e+00]
[ 2.34718235e+00 -1.09297809e+00]
[-9.64145695e-01 -4.29624864e-01]
[ 1.45567095e+00  1.34225380e-01]
[-9.64145695e-01 -1.35831938e+00]
[-7.09428153e-01 -9.27139784e-01]
[-3.27351841e-01  7.64410946e-01]
[-7.26342990e-02  6.31740300e-01]

```



```
[-4.54710612e-01  9.30249252e-01]
[ 5.64159555e-01  3.99566671e-01]
[-1.99993070e-01  2.88714127e+00]
[ 1.96510603e+00 -1.25881640e+00]
[ 5.47244718e-02  4.32734332e-01]
[-5.82069382e-01 -1.52415769e+00]
```

```
[17]: # Training the Logistic Regression model on the Training set
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0)
classifier.fit(X_train, y_train)
```

```
[17]: LogisticRegression(random_state=0)
```

```
[18]: # Predicting a new result
print(classifier.predict(sc.transform([[30,87000]])))
```

[0]

```
[19]: # Predicting the Test set results
y_pred = classifier.predict(X_test)
print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.
    ↪ reshape(len(y_test),1)),1))
```

$$\begin{bmatrix} [0 & 0] \\ [0 & 0] \\ [0 & 1] \\ [0 & 0] \\ [0 & 0] \\ [0 & 0] \\ [1 & 1] \\ [0 & 0] \\ [0 & 0] \\ [0 & 0] \\ [0 & 1] \\ [0 & 0] \\ [0 & 0] \\ [0 & 1] \\ [0 & 0] \\ [0 & 0] \\ [0 & 0] \\ [0 & 1] \\ [0 & 0] \\ [0 & 0] \\ [0 & 0] \\ [0 & 0] \\ [0 & 0] \\ [0 & 0] \end{bmatrix}$$

```

[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[1 1]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 0]
[0 1]
[0 1]
[0 0]
[0 0]]

```

```

[20]: # Making the Confusion Matrix
from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)

```

```

[[35  0]
 [ 6  2]]

```

[20]: 0.8604651162790697

```

[42]: # Visualising the Training set results
from matplotlib.colors import ListedColormap
import numpy as np
import matplotlib.pyplot as plt

# Inverse transform to get the original scale
X_set, y_set = sc.inverse_transform(X_train), y_train

# Create the meshgrid with a larger step size
step_size_age = 2 # Increased step size for age
step_size_salary = 500 # Increased step size for estimated salary

# Generate the meshgrid for plotting
X1, X2 = np.meshgrid(
    np.arange(start=X_set[:, 0].min() - 10, stop=X_set[:, 0].max() + 10,
↪step=step_size_age),
    np.arange(start=X_set[:, 1].min() - 1000, stop=X_set[:, 1].max() + 1000,
↪step=step_size_salary)

```

```

)

# Predict the classifier output and reshape
Z = classifier.predict(sc.transform(np.array([X1.ravel(), X2.ravel()]).T)).
    ↪ reshape(X1.shape)

# Plotting
plt.contourf(X1, X2, Z, alpha=0.75, cmap=ListedColormap([(1, 0, 0), (0, 1, 0),
    ↪ (0, 0, 1)])) # Using RGB tuples for colors
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())

# Scatter plot the training points
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c=ListedColormap([(1, 0, 0), (0, 1, 0)])(i), label=j) # Same
    ↪ color scheme


    # Alternatively, you can use:
    # c=['red', 'green'][i]

plt.title('Logistic Regression (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()

```

C:\Users\ddaya\AppData\Local\Temp\ipykernel_15000\4181641556.py:29: UserWarning:
 c argument looks like a single numeric RGB or RGBA sequence, which should be
 avoided as value-mapping will have precedence in case its length matches with
 x & *y*. Please use the *color* keyword-argument or provide a 2D array with a
 single row if you intend to specify the same RGB or RGBA value for all points.

```
plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
```



```
[44]: # Visualising the Training set results
from matplotlib.colors import ListedColormap
import numpy as np
import matplotlib.pyplot as plt

# Inverse transform to get the original scale
X_set, y_set = sc.inverse_transform(X_test), y_test

# Create the meshgrid with a larger step size
step_size_age = 2 # Increased step size for age
step_size_salary = 500 # Increased step size for estimated salary

# Generate the meshgrid for plotting
X1, X2 = np.meshgrid(
    np.arange(start=X_set[:, 0].min() - 10, stop=X_set[:, 0].max() + 10,
    ↪step=step_size_age),
    np.arange(start=X_set[:, 1].min() - 1000, stop=X_set[:, 1].max() + 1000,
    ↪step=step_size_salary)
)
```

```

# Predict the classifier output and reshape
Z = classifier.predict(sc.transform(np.array([X1.ravel(), X2.ravel()]).T)).
    ↪reshape(X1.shape)

# Plotting
plt.contourf(X1, X2, Z, alpha=0.75, cmap=ListedColormap([(1, 0, 0), (0, 1, 0)],
    ↪0))) # Using RGB tuples for colors
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())

# Scatter plot the training points
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c=ListedColormap([(1, 0, 0), (0, 1, 0)])(i), label=j) # Same
    ↪color scheme
    # Alternatively, you can use:
    # c=['red', 'green'][i]

plt.title('Logistic Regression (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()

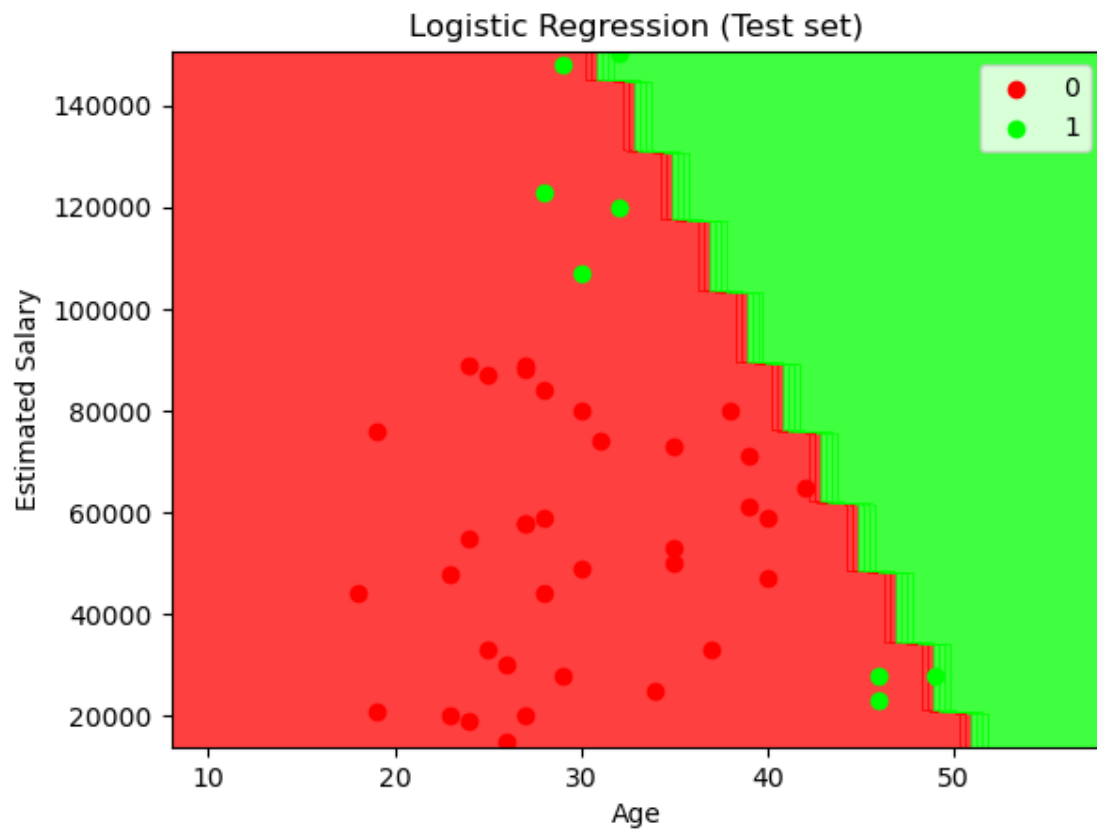
```

C:\Users\ddaya\AppData\Local\Temp\ipykernel_15000\2172703052.py:29: UserWarning:
 c argument looks like a single numeric RGB or RGBA sequence, which should be
 avoided as value-mapping will have precedence in case its length matches with
 x & *y*. Please use the *color* keyword-argument or provide a 2D array with a
 single row if you intend to specify the same RGB or RGBA value for all points.

```

plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],

```



[]: